



Predictive Analytics for Sales Forecasting

Prepared By: Priyanshu Singh
Master of Science in **Data Science**

Date: 19th May 2026

ABSTRACT

This project is a comprehensive implementation of predictive analytics solutions for sales forecasting for the financial and insurance verticals. This study is called Predictive Analytics for Sales Forecasting in the Finance Sector, and the data set it uses includes historical information from insurance agencies from 2005 to 2014 (213,328 records from 1,623 insurance agencies in 6 states of the United States).

The goal of this project is to build and test predictive models which would be able to predict the written premium accurately from past financial metrics such as loss ratio, retention ratio, new business premium, and incurred losses. The methods that were employed are Linear Regression (baseline), Random Forest Regressor, and Facebook Prophet for time-series forecasting, all three of which were implemented and compared.

This project uses Python as the main programming language, and some of the libraries utilized are Pandas, NumPy, Scikit-learn, and Prophet for data processing and machine learning respectively. Structured data was queried using SQL (SQLite) whereas Power BI was used to create an interactive 3-page dashboard, offering a descriptive and predictive analytical insight.

The best results were obtained by the Random Forest Regressor on log-transformed premium values with an R^2 score of 0.848, a Mean Absolute Error (MAE) of 0.505 and an RMSE of 0.858. The two best indicators of total premium volume are incurring losses and new business written premium, according to key findings. Commercial lines grew 114% over the last decade, and Ohio made up 58% of the total written premium portfolio.

The project builds the existing knowledge in the field of predictive analytics in finance and offers valuable information for data driven recommendations to enhance the accuracy of sales forecasting in the insurance industry.

TABLE OF CONTENTS

Chapter 1: Introduction	1
1.1 Background of the Project	1
1.2 Problem Statement	2
1.3 Objectives of the Project	2
1.4 Scope of the Project	3
1.5 Existing System Overview	3
1.6 Proposed System Overview	4
1.7 Technologies Used	5
Chapter 2: Literature Review	6
2.1 Introduction to Predictive Analytics in Finance	6
2.2 Related Studies Review	7
2.3 Comparative Analysis of Techniques	10
2.4 Software Development Methodology	10
2.5 Relevant Frameworks and Libraries	11
2.6 Research Gap	11
Chapter 3: System Analysis	13
3.1 Functional Requirements	13
3.2 Non-Functional Requirements	14
3.3 Feasibility Study	14
3.4 System Architecture	15
3.5 Data Flow Diagrams	16
3.6 Use Case Diagrams	18
Chapter 4: System Design	20
4.1 UML Diagrams	20
4.2 ER Diagram	24
4.3 Database Schema	25
Chapter 5: System Implementation	26
5.1 Development Environment	26
5.2 Tools and Technologies	26
5.3 Dataset Description and EDA	26

5.4 Data Preprocessing	31
5.5 Predictive Model Development	33
5.6 SQL Implementation	36
5.7 Power BI Dashboard	40
Chapter 6: Testing	42
6.1 Testing Strategy	42
6.2 Unit Testing	42
6.3 Integration Testing	43
6.4 System Testing	44
Chapter 7: Results and Discussion	45
7.1 EDA Results	45
7.2 Model Performance Results	45
7.3 Dashboard Results	47
7.4 Comparison with Existing Systems	49
Chapter 8: Conclusion and Future Enhancements	50
8.1 Conclusion	50
8.2 Future Scope	51
Chapter 9: References	52
Chapter 10: Appendices	55

LIST OF FIGURES

Figure No.	Figure Title	Page No.
Figure 5.1	Total Written Premium by Year (2005–2014)	
Figure 5.2	Written Premium by Product Line (2005–2014)	
Figure 5.3	Total Written Premium by State	
Figure 5.4	Average Loss Ratio vs Retention Ratio (2005–2014)	
Figure 5.5	Written Premium Distribution (Raw vs Log-Transformed)	
Figure 5.6	Correlation Heatmap — Key Financial Indicators	
Figure 5.7	Random Forest — Predicted vs Actual (Log Scale)	
Figure 5.8	Facebook Prophet — Sales Forecast (2005–2014)	
Figure 5.9	Predictive Model Performance Comparison	
Figure 5.10	Random Forest — Feature Importance	
Figure 7.1	Power BI — Sales Overview Dashboard	
Figure 7.2	Power BI — Financial Health Dashboard	
Figure 7.3	Power BI — Predictive Analytics Dashboard	

LIST OF TABLES

Table No.	Table Title	Page No.
Table 1.1	Technologies Used in the Project	
Table 3.1	Functional Requirements	
Table 3.2	Non-Functional Requirements	
Table 3.3	Feasibility Analysis Summary	
Table 4.1	Database Schema - insurance sales Table	
Table 5.1	Dataset Overview	
Table 5.2	Key Financial Columns and Definitions	
Table 5.3	Preprocessing Steps Summary	
Table 5.4	Feature Engineering Summary	
Table 6.1	Unit Test Cases	
Table 6.2	Integration Test Cases	
Table 6.3	System Test Cases	
Table 7.1	Model Performance Comparison	
Table 7.2	SQL Query Results Summary	

LIST OF ABBREVIATIONS

Abbreviation	Full Form
ML	Machine Learning
EDA	Exploratory Data Analysis
MAE	Mean Absolute Error
RMSE	Root Mean Square Error
R ²	R-Squared (Coefficient of Determination)
RF	Random Forest
LR	Linear Regression
CL	Commercial Lines
PL	Personal Lines
SQL	Structured Query Language
BI	Business Intelligence
DFD	Data Flow Diagram
UML	Unified Modelling Language
ER	Entity Relationship
API	Application Programming Interface
CSV	Comma-Separated Values
KPI	Key Performance Indicator
YoY	Year-over-Year
UI	User Interface

CHAPTER 1: INTRODUCTION

1.1 Background of the Project

The insurance and financial services sector create huge transactional data on a daily basis. Insurance agencies monitor written premiums, policy renewals, claims, loss ratios and new business, as well as thousands of client accounts. Despite the wealth of historical information available, a lot of financial organizations are still using manual approaches, gut instinct or cruder statistical models to forecast sales success in the future. Predictive Analytics is the predictive modelling process to identify future events from past trends using statistical algorithms, machine learning models, and data mining techniques. As it relates to insurance sales, predictive analytics can revolutionize the way that agencies allocate resources, target their sales goals and find markets where growth is going to occur before it does. Being able to predict written premium volumes one to three years ahead gives insurance agencies an edge in the competitive arena. The financial services industry - including banking, insurance, investment management and financial advisory - is one of the most data-intensive sectors in the global economy. Recent studies indicate that companies with sophisticated forecasting methods using advanced analytics can outperform companies with traditional forecasting methods in terms of forecast accuracy and operational efficiency. Predictive analytics is now possible for every organization, thanks to the rise of cloud computing, open-source machine learning libraries, and business intelligence platforms. This project is targeted specifically at the insurance sub-sector of the finance sector with a real-world dataset of agency performance measures covering 11 years (2005 - 2015) of data. The data includes 1,623 independent insurance agencies in six states in the United States, offering a variety of commercial lines (CL) and personal lines (PL) insurance products. The project investigates how to use several machine learning approaches to predict written premium sales and compare the models' performances using conventional statistical measures.

1.2 Problem Statement

One of the major challenges for insurance agencies and financial institutions in the regional insurance market is the lack of ability to correctly estimate future volumes of written premiums from historical volumes. Traditional forecasting approaches focus primarily on year-over-year growth projections and/or a manual examination of the past, without taking into consideration the various factors that influence financial indicators like loss ratios, retention rates, acquisition of new business, market specific factors, etc. The lack of a systematic, data-driven forecasting framework means there are several operational issues. Agencies cannot allocate their resources and establish realistic sales goals without first being able to do the above. Second, without knowledge of expected claim trends and volumes, underwriters are unable to sufficiently price the risk. Third, management does not have visibility of what product lines, geographic areas and agency segments are likely to be more successful or less successful in the next several years. To meet these challenges, this project builds and tests predictive models that will be able to estimate insurance sales (written premium) based on financial performance indicators from past years and compare their performance. The proposed system is designed to fill the gap between the amount of available data and the ability to take action to achieve the desired forecasting.

1.3 Objectives of the Project

To conduct a comprehensive literature review and understand the significance of predictive analytics in sales forecasting in the finance industry. Analyses past sales data and financial metrics to create predictive models for insurance premium forecasting. Compare the accuracy and effectiveness of multiple predictive models such as Linear Regression model, Random Forest Regressor model and Facebook Prophet model using standard evaluation metrics (MAE, RMSE, R^2). To learn about various predictive analytics methods to forecast sales and to understand the best model for the insurance industry. To make possible recommendations to enhance sales forecasting through predictive analytics in the finance and insurance sectors.

1.4 Scope of the Project

This project's scope includes the entire data pipeline from data acquisition to business recommendations. The project includes the following activities:

- A project involving Exploratory Data Analysis (EDA) for the analysis of 213,328 records of insurance agency performance over the time period 2005-2014.
- Data preprocessing (sentinel values of treatment, feature engineering, log transformation of the target variable, etc.), and train test splitting. Design and testing machine learning sales forecasting models.
- Utilization of Structured Query Language (SQL) for aggregation of agency performance and financial health assessment using the SQLite language.
- Creating an Interactive Dashboard with Microsoft Power BI with 3 Analytical Pages. Structured academic report including documentation of findings, methodology and recommendations.

This project is restricted to the insurance industry and uses the publicly available Kaggle data. It does not involve real-time data streaming, proprietary insurance systems integration, or deployment in a production environment.

1.5 Existing System Overview

The three main types of sales forecasting currently used in insurance and finance are, there are three major methods of current sales forecasting used in insurance and finance: The first are the manual projections, in which sales managers estimate future performance based upon past averages and growth rate assumptions. They are all subjective and open to human bias and can be prone to inaccurate forecasts which fail to take into account underlying market dynamics. The other types of simple models are basic statistical models like moving averages and simple linear regression. These models are more objective but cannot incorporate non-linearities and complex interactions among financial variables. They also have to deal with the variable and skew nature of insurance premiums. The third category includes commercial enterprise resource planning (ERP) systems and actuarial software with a certain forecasting component, but are costly, specialized software, and difficult to modify exploratory analytical workflows. These

systems tend to be opaque and offer no detail at the feature level of what is driving sales. Current systems have a few drawbacks: They have no machine learning ability, they don't provide quantitative feature importance, they don't provide any time-series aware forecasting, and they don't allow for predictive insights to be visually communicated to business users.

1.6 Proposed System Overview

The suggested system is an insurance sales forecasting application that utilizes end-to-end predictive analytics pipeline. The system takes historical performance data of agency as input and outputs forecast results with a series of well-defined analytical phases.

It starts with data ingestion and exploratory data analysis, followed by systematic pre-processing to tackle data quality challenges such as sentinel value treatment, feature engineering, and target variable normalization. Three predictive models are then trained and evaluated: Linear Regression model as a baseline model, a Random Forest Regressor model to capture the non-linear patterns, and a Facebook Prophet model to forecast the time-series specific patterns.

The system outputs performance metrics of the models (MAE, RMSE, R2), feature importance rankings, and comparison charts of forecast vs actual. These are fed into a three-page Power BI dashboard that gives stakeholders descriptive, financial health and predictive analytics views. An SQL layer allows for structured queries to be used for aggregate reporting.

The proposed system is more advantageous than the current ones, as it integrates the transparency of machine learning (feature importance), the time-series aspect (Prophet), and the interactive visualization (Power BI) in a single analytical system, which is reproducible, developed entirely using open source tools.

1.7 Technologies Used

Technology	Version/Tool	Purpose
Python	3.13	Data analysis, ML modeling, preprocessing
Pandas	2.x	Data manipulation and analysis
NumPy	1.x	Numerical computation
Scikit-learn	1.x	Linear Regression, Random Forest, metrics
Facebook Prophet	1.x	Time-series forecasting
Matplotlib / Seaborn	3.x	Data visualization
SQLite (Python)	Built-in	Database creation and SQL queries
Microsoft Power BI	Desktop	Interactive dashboard
Jupyter Notebook	VS Code	Development environment
Git / GitHub	Latest	Version control and portfolio hosting
Kaggle	Platform	Dataset source

CHAPTER 2: LITERATURE REVIEW

2.1 Introduction to Predictive Analytics in Finance

Predictive analytics is one of the advanced analytics subfields, which involves forecasting future events or outcomes based on historical data, statistical algorithms, and machine learning techniques. Predictive analytics has proved to be a key ability in the finance industry, where businesses are looking to make data-informed decisions to gain a competitive edge. The financial services sector spans banking, insurance, investment management, and lending and creates huge amounts of transactional and behavioural data that can prove to be very insightful when analysed correctly and in good time, allowing for predictions to be made.

In recent years, the use of predictive analytics in financial sales forecasting has grown significantly, as financial institutions have realized that conventional methods of forecasting are not effective in today's complex market conditions, and as more data has become available in the historical record, open source machine learning frameworks have proliferated. Machine learning-based forecasting has been shown to consistently outperform traditional statistical forecasting techniques in terms of forecast error.

Predictive Analytics is used in insurance for several purposes: premium pricing, predicting claims, forecasting customer retention and managing the pipeline of new business. Being able to predict written premium volume at the agency level is especially useful for regional insurance groups that have multiple lines of business and multiple independent agents in various states.

2.2 Review of Related Studies

Study 1: Prediction for Insurance Premiums Based on Random Forest and Multiple Linear Regression

A study published on researchgate (2023) compared the performance of the Random Forest model to the Multiple Linear Regression model in predicting insurance benefits and premium amounts. The researchers tested both models on an actual insurance dataset and demonstrated the superiority of the Random Forest model over Linear Regression in prediction accuracy from the dataset. The study also suggested that although the accuracy of Random Forest was higher, it was not easily interpretable as was Linear Regression where the relationship between the input variables and the premium outcomes can be easily explained.

The authors found that both models are useful for their respective organizational purposes: Random Forest is better for the accuracy of the task, whereas Linear Regression is suitable for people who want to know more about forecasting and need to understand the process. The approach taken in this project to implement both models and evaluate the trade-offs fits well with this finding. The same trend has been observed in our results, which shows that the R^2 value of the Random Forest is higher than the Linear Regression's on the insurance agency dataset, with $R^2 = 0.848$ and $R^2 = 0.309$, respectively.

Study 2: Random Forest Regression with Hyper Parameter Tuning for Medical Insurance Premium Prediction

Prakash et al. (2022) conducted a research paper that showed the effectiveness of Random Forest Regression (with hyperparameter tuning) for medical insurance premium prediction in the International Journal of Health Sciences. The study was applied to a healthcare insurance dataset, showing that the model could be optimized with good hyperparameter values, which would lead to better performance. Data preprocessing, specifically feature selection, as well as outlier detection, was found to be essential for obtaining accurate prediction outcomes.

There are two major aspects of the paper that are relevant to this project. First, it confirms the selection of the Random Forest model as the main predictive model to forecast insurance premiums. Secondly it highlights the significance of the preprocessing pipeline used in this project, specifically the treatment of the heavily skewed target variable and the sentinel values (99999 placeholders).

Study 3: A Study of Impact and Applications of Predictive Analytics in Sales Forecasting

Kumar (2023) conducted an extensive review of the International Journal of Research in Applied Science and Engineering Technology on the various applications and applications of predictive analytics in various sectors including finance, retail, and healthcare. The study revealed that AI-driven predictive systems can crunch through vast amounts of historical and pipeline data and deliver much more accurate sales forecasts than human analysis. The organizations that used predictive sales forecasting saw an increase in revenue predictability, in efficient use of resources, and capability for better strategic planning.

A crucial aspect of this paper for this project is the fact that it provides the theoretical foundation to use machine learning for sales forecasting in the finance industry. The results validate the project's main thesis of generating reliable forward-looking premium forecasts through the analysis of historical agency performance with the right ML models. The paper also states that data quality and feature engineering are two challenges that can become a success factor, and therefore were taken into account in this project when making choices among the different pre-processing options.

Study 4: AI Driven Predictive Analytics for Financial Forecasting.

On the other hand, the IEEE conference paper (2024) explores the application of AI predictive analytics for financial forecasting in risk management and business planning. The study evaluated the performance of AI technologies for predicting financial results and their potential for real-time reporting in financial institutions. The study revealed that companies that adopted AI-driven forecasting systems outperformed those employing

traditional statistical methods when it came to the accuracy of their financial performance forecasts.

The paper's main finding is that predictive analytics can be a game-changer in financial risk management. It determined that predictive models with multiple financial indicators outperformed models with single variable forecasting, which directly supports the multi-feature modeling approach taken for this project. The features used in our model (loss ratio, retention ratio, new business premium) are consistent with the multi-indicator strategy suggested by this research.

Study 5: Predictive Analytics in Financial Management: Enhancing Decision-Making and Risk Management

A systematic review of the use of predictive analytics to improve decision-making and risk management in financial services organisations was published on ResearchGate (2024). The review compiled results from diverse financial management studies and concluded that predictive analytics in financial management drive tangible benefits in terms of enhanced predictive accuracy, risk measurement, and operational efficiency. Specifically, the review referenced a survey conducted by the World Economic Forum (2023) which revealed that 54% of the financial sector firms faced challenges in recruiting individuals with the appropriate skills in analytics.

The literature review revealed a clear gap in the literature with regards to using predictive analytics in insurance sales forecasting at the agency level over multiple years—although there is a lot of literature in the e-commerce and banking sectors to support the use of predictive analytics in cost prediction for medical insurance and banking, respectively, there is limited published research that focuses on multi-year agency-level insurance sales forecasting using ensemble methods. This gap directly underlies the direction of the research of this project.

2.3 Comparative Analysis of Predictive Techniques

Technique	Accuracy	Interpretability	Best Use Case
Linear Regression	Low-Medium	High	Baseline, transparent models
Random Forest	High	Medium	Complex non-linear patterns
Facebook Prophet	Medium	High	Seasonal time-series
XGBoost	Very High	Low	Tabular data competitions
LSTM/Deep Learning	High	Very Low	Large sequential datasets

Considering the nature of the insurance agency performance data and the comparative analysis, the Random Forest Regressor model was chosen as the primary model because of its ability to provide a good balance of accuracy and interpretability. The interpretable baseline was Linear Regression and Facebook Prophet was added because of its time-series awareness.

2.4 Software Development Methodology

The project is undertaken in accordance with the methodology of Cross-Industry Standard Process for Data Mining (CRISP-DM) which is an industry standard method for data science and data analytics projects. The CRISP-DM is made up of six phases: Business Understanding, Data Understanding, Data Preparation, Modeling, Evaluation, Deployment.

The CRISP-DM methodology was chosen over the traditional software development models (Waterfall, Agile) as the methodology is specifically designed for analytical and machine learning projects that are primarily looking to deliver a predictive model which is different from a software application. It makes it possible to go back and refine the data preparation and modelling steps as a result of the evaluations performed, as necessary in this project when dealing with data quality concerns, e.g. sentinel values or distribution skewness.

2.5 Relevant Frameworks and Libraries

The scikit-learn library is an open-source Python machine learning library that includes implementations of Linear Regression, Random Forest Regressor, and evaluation metrics such as MAE, RMSE, and R^2 . It was chosen because of its API's well documented use, reproducibility, and adoption in the data science community.

Facebook Prophet is a time-series forecasting open-source library, created by the Core Data Science team at Facebook. This method is designed to process datasets that have strong seasonal characteristics and multiple seasonalities, and has a good performance on the time-series with missing data and outliers. The additive model structure of Prophet is especially well suited for forecasting annual insurance premiums.

Microsoft Power BI is a business analytics service that offers interactive data visualization and business intelligence. It was chosen for dashboard development for a number of reasons including built-in support for DAX calculated columns, slicer-based filtering and multi-page report layouts.

SQLite is a small, embedded, serverless, relational database engine that comes with Python. It was employed to develop a structured database from the cleaned data and to run the SQL aggregation queries to analyze agency performance.

2.6 Research Gap

A review of literature clearly indicates that this project fills a gap in the literature. Although there is a considerable amount of research on predictive analytics for fraud detection in banking, predictive credit risk assessment and medical insurance cost prediction, there is limited published research focusing on multi-year agency-level insurance sales forecasting that integrates ensemble machine learning techniques with time-series forecasting and interactive business intelligence visualization within a single analytical framework.

In particular, some previous works focus on the performance of a single model only, or on forecasting in one domain only. The existing studies reviewed do not integrate the three models - Random Forest, Linear Regression, and Prophet - together and compare their

performance on insurance agency performance data, using explicit financial metrics like loss ratio, retention ratio, or new business premium as predictive features.

To fill this gap, this project will show a comparative performance study of three forecasting techniques on a real-world insurance dataset and discuss the results in a multi-page Power BI dashboard - all in one end-to-end reproducible, open-source analytical pipeline.

CHAPTER 3: SYSTEM ANALYSIS

3.1 Functional Requirements

FR No.	Requirement	Description
FR-01	Data Ingestion	The system shall load CSV data into Python DataFrame and SQLite database.
FR-02	Data Exploration	The system shall generate summary statistics and EDA visualizations.
FR-03	Data Preprocessing	The system shall handle sentinel values, encode categoricals, and apply log transformation.
FR-04	Model Training	The system shall train Linear Regression, Random Forest, and Prophet models.
FR-05	Model Evaluation	The system shall compute MAE, RMSE, and R^2 for each model.
FR-06	Visualization	The system shall generate forecast vs. actual charts and feature importance plots.
FR-07	SQL Querying	System shall execute aggregation queries on the SQLite database.
FR-08	Dashboard	The system shall display results in an interactive Power BI dashboard.

3.2 Non-Functional Requirements

NFR No.	Category	Requirement
NFR-01	Performance	Model training shall complete within 5 minutes on standard hardware.
NFR-02	Reliability	Preprocessing pipeline shall handle missing/sentinel values without data loss.
NFR-03	Reproducibility	All notebooks shall use fixed random seeds (random_state=42).
NFR-04	Scalability	The system shall handle datasets up to 500,000 records.
NFR-05	Usability	Dashboard shall be navigable by non-technical stakeholders.
NFR-06	Portability	The project shall run on Windows, Mac, and Linux environments.

3.3 Feasibility Study

Technical Feasibility

The project can be technically implemented with easily available open source tools. The data processing and model development are all covered by Python 3.13 and its libraries Pandas, Scikit-learn and Prophet. Power BI Desktop is a free download from Microsoft. SQLite comes built into Python and doesn't need to be installed. The dataset (213,328 records, 49 columns) is not too complex to process with the laptop hardware.

Economic Feasibility

There are low cost implications of the project. The software tools used (Python, Scikit-learn, Prophet, SQLite, Jupyter Notebook, GitHub) are all open source and free to use. Power BI Desktop is available for free for individuals and for educational purposes. This

data was obtained free of charge from Kaggle. No cloud computing or infrastructure costs since all processing is done locally.

Operational Feasibility

The project is within the scope of the academic year. The CRISP-DM methodology is structured and can be done in a 14-day sprint. This report, along with the other outputs such as Jupyter Notebooks, SQL query, and the dashboard (Power BI), are well-documented and can be maintained and extended by future researchers.

3.4 System Architecture

The system is architected as four layers: Data Layer, Processing Layer, Analytics Layer and Presentation Layer.

The Data Layer is the raw CSV data that is located in the data/raw/ directory as well as the processed data in the data/processed/ directory, and the SQLite database

(insurance.db). The Processing Layer has three jupyter notebooks:

01_data_exploration.ipynb for data exploration, 02_preprocessing_modeling.ipynb for data preprocessing and modeling, and 03_sql_queries.ipynb for SQL-based analysis. The Analytics Layer is the trained models (Linear Regression, Random Forest, Prophet) and their prediction output files. The Power BI dashboard with its three pages and the project report are part of the Presentation Layer.




Figure 3.1: System Architecture Diagram

3.5 Data Flow Diagrams

DFD Level 0 (Context Diagram)

The Level 0 DFD displays the system as one process with the external entities and data flows. External users are: Data Analyst (user), Kaggle Dataset (data source), and

Business Stakeholder (report consumer). The key data flows include: Raw CSV data coming in from Kaggle to the system; Processed forecasts and insights from the system to the Data Analyst; and Dashboard visualizations to the Business Stakeholder.




Figure 3.2: DFD Level 0 - Context Diagram

DFD Level 1

The Level 1 DFD further breaks down the system into its key processes:

1. Data Ingestion and Exploration,
2. Data Preprocessing,
3. Model Training and Evaluation,
4. SQL Analysis
5. Dashboard Generation. Data stores (raw data store, processed data store, model output store) and data flows (data transformation) link each process.




Figure 3.DFD Level 1

3.6 Use Case Diagrams

The key players in the system are Data Analyst and Business Stakeholder. The Data Analyst is the person who communicates with the system in order to load, preprocess, train the model, query the SQL, and create the dashboard. The Business Stakeholder can use the system to view forecasts, filter by year/product line and export reports.




Figure 3.4: Use Case Diagram

CHAPTER 4: SYSTEM DESIGN

4.1 UML Diagrams

4.1.1 Use Case Diagram

The following Use Case Diagram illustrates the interactions between the Data Analyst and the various Use Cases. The following Use Case Diagram shows the interactions between the Data Analyst and the Use Cases: Load Data, Preprocess, Train Model, Evaluate, Query SQL, Build Dashboard.



Use Case Diagram - Actor: Data Analyst, Use Cases: Load Data, Preprocess, Train Model, Evaluate, Query SQL, Build Dashboard

4.1.2 Class Diagram

This class diagram is used to show the main data structures and their relationships in this project. The main classes are: DataLoader for loading CSV and database,

DataPreprocessor for cleaning, encoding and feature engineering, ModelTrainer for encapsulating training logic over the three models, ModelEvaluator for computing MAE, RMSE and R^2 , and DashboardExporter for CSV exporting for Power BI.




Figure 4.2: Class Diagram

4.1.3 Sequence Diagram

The sequence diagram depicts the interactions between the Data Analyst, Python notebooks, and the database in the model training workflow. The pipeline starts by loading the CSV, passing it to the pre-processing module, continuing with training, evaluation, and finally the prediction module is passed to the CSV for the consumption of

Power BI.




Figure 4.3: Sequence Diagram

4.1.4 Activity Diagram

The activity diagram depicts the entire process, from data loading to the generation of a report, and also the handling of missing data and the selection of the optimal model with the maximum R^2 value.




Figure 4.4: Activity Diagram

4.2 ER Diagram

The Entity-Relationship Diagram represents the structure of the insurance_sales table in the SQLite database. The main table is insurance_sales which contains the attributes AGENCY_ID (which is the primary key), STAT_PROFILE_DATE_YEAR, PROD_LINE_ENC, STATE_ENC, WRTN_PREM_AMT, LOSS_RATIO, and RETENTION_RATIO, and some other financial attributes.




Figure 4.5: ER Diagram

4.3 Database Schema

Column Name	Data Type	Description
AGENCY_ID	INTEGER	Unique identifier for each insurance agency
STAT_PROFILE_DATE_YEAR	INTEGER	Year of the record (2005–2014)
POLY_INFORCE_QTY	INTEGER	Number of policies in force
NB_WRTN_PREM_AMT	REAL	New business written premium amount (USD)
PRD_INCRD_LOSSES_AMT	REAL	Period incurred losses amount (USD)
RETENTION_RATIO	REAL	Customer 2retention ratio (0–1)
LOSS_RATIO	REAL	Loss ratio (losses/premium)
GROWTH_RATE_3YR	REAL	3-year premium growth rate
ACTIVE_PRODUCERS	INTEGER	Number of active sales producers
AGENCY_AGE	INTEGER	Age of agency (engineered feature)
PREMIUM_GROWTH_YOY	REAL	Year-over-year premium growth (engineered)
PROD_LINE_ENC	INTEGER	Product line encoded (0=CL, 1=PL)
STATE_ENC	INTEGER	State encoded (0–5)
MONTHS	INTEGER	Number of months covered in record
LOG_WRTN_PREM_AMT	REAL	Log-transformed written premium (target)
WRTN_PREM_AMT	REAL	Written premium amount in USD (raw target)

CHAPTER 5: SYSTEM IMPLEMENTATION

5.1 Development Environment

- Operating System: Windows 11
- Programming Language: Python 3.13
- IDE: Jupyter Notebook
- Version Control: Git 2.x with GitHub
- Database: SQLite (via Python sqlite3 module)
- Dashboard: Microsoft Power BI Desktop
- Hardware: Intel Core i5 processor, 8GB+ RAM, 512GB SSD

5.2 Tools and Technologies

Library/Tool	Version	Purpose
pandas	2.x	Data loading, manipulation, aggregation
numpy	1.x	Numerical operations, array processing
matplotlib	3.x	Static chart generation
seaborn	0.x	Statistical visualization (heatmap)
scikit-learn	1.x	Linear Regression, Random Forest, metrics
prophet	1.x	Time-series forecasting
sqlite3	Built-in	SQL database operations
Power BI Desktop	Latest	Interactive dashboard

5.3 Dataset Description and EDA

In this project, I'll be using the Insurance Agency Performance Dataset from Kaggle (moneystore/agencyperformance). It includes performance data for 1,623 independent insurance agencies, a group of 10 states with insurance companies that provide property, casualty, life and brokerage products in Ohio (OH), Pennsylvania (PA), Kentucky (KY), Indiana (IN), West Virginia (WV) and Michigan (MI).

Attribute	Value
Total Records	213,328
Total Features	49
Time Period	2005–2015
Unique Agencies	1,623
States Covered	6 (OH, PA, KY, IN, WV, MI)
Product Lines	Commercial (CL) and Personal (PL)
Null Values	Zero (complete dataset)
Duplicate Records	Zero

5.3.1 Annual Premium Trend Analysis

The total written premium for 2005-2014 is illustrated in Figure 5.1. The overall growth of the period is found from 2005 (\$274.8M) to 2006 (\$413.4M) with an annual growth of 50.5%. This was followed by a stable plateau from 2006 to 2014 during which total premiums were consistently in the \$400-\$440M range. The portfolio was largest in 2014 at \$437.5M, signifying that the company has a mature and stable insurance portfolio and has a stable customer base.




Figure 5.1: Annual Premium Trend

5.3.2 Premium by Product Line

The disparity in trends of Commercial Lines (CL) and Personal Lines (PL) can be seen in the figure 5.2 below. Throughout the period, Personal Lines premium volume consistently was higher than Commercial Lines premium volume. A significant change in strategy is being seen however – Commercial Lines has increased from \$95.5M in 2005 to \$204.4M in 2014 with 114% growth, while Personal Lines has decreased from \$179.4M in 2005 to \$233.1M in 2014, with a downtrend after 2006. This indicates a slow transition towards a more commercial insurance product strategy.




Figure 5.2: Premium by Product Line

5.3.3 Geographic Premium Distribution

The total written premiums by state for the entire study period are shown in Figure 5.3. The Ohio (OH) state portfolio makes up about 58% of the total written premium, led by Ohio. Pennsylvania (\$554.9M), Kentucky (\$491.8M) and Indiana (\$454.7M) are next with significantly lower volumes. Michigan is the lowest contributor with \$45.7M. This geographical focus implies to the portfolio risk and indicates the expansion strategy needed in the lower penetration states.




Figure 5.3: Premium by State

5.3.4 Loss Ratio and Retention Analysis

The graph (Figure 5.4) shows the average loss ratio and the average retention ratio over time, plotted on a dual axis. Average retention ratio for the portfolio is 88.75%, a sign of high levels of customer affinity. However, the loss ratio is very volatile, with an unusual high peak in 2012 which could be a sign of increased claims period. Loss ratios were pushed back to more normalized levels in 2013-2014 following the increase in 2012 which indicated that underwriting corrections were successful.




Figure 5.4: Loss vs retention

5.3.5 Premium Distribution Analysis

As shown in figure 5.5, the written premium distribution is very skewed towards the right. The raw distribution reveals that most of the records have very little premium value (close to \$0), and high-performing agencies earn premiums of more than \$1M. The distribution of the target variable is close to normal after log-transforming, confirming the use of log1p transformation prior to model training. This transformation greatly enhanced the performance of the models.




Figure 5.5: Premium Distribution

5.3.6 Correlation Analysis

The correlation heatmap for key financial indicators is given in Figure 5.6. The most important result is that there is almost perfect correlation ($r = 1.00$) between PRD_ERND_PREM_AMT and WRTN_PREM_AMT, data leakage (both measure the same underlying construct). The column was not included in the features. The number of active policies (POLY_INFORCE_QTY) has the most strong positive correlation with the target variable ($r = 0.81$), indicating that the most significant predictor of premium volume is the number of active policies.




Figure 5.6: Correlation Heatmap

5.4 Data Preprocessing

5.4.1 Data Filtering

The data from 2015 was not used in the analysis. This is because I didn't have all the months for that year. If we used it it would not be fair to the years and it would mess up the trends we are trying to find. We also did not use records that had values, for the written premium. These are cases where policies were cancelled or reversed. They are not sales. After I removed all these records we had 147,760 records left. I started with 213,328 records.

5.4.2 Sentinel Value Treatment

The dataset had some numbers (99999) that meant the data was missing. These numbers were used in financial columns.

- They were found in GROWTH_RATE_3YR (50,546 times)
- LOSS_RATIO_3YR (16,209 times)

- RETENTION_RATIO (88,984 times).
- ACTIVE_PRODUCERS (several times).

These special numbers were replaced with not a number" value (NaN).

The missing values were then filled in with the value of each column.

This method works with financial data that has extreme values.

The dataset now has GROWTH_RATE_3YR, LOSS_RATIO_3YR, RETENTION_RATIO. With accurate numbers. The financial ratio columns now have data. The numbers were fixed to make sure they are useful for analysis. The sentinel values, in GROWTH_RATE_3YR LOSS_RATIO_3YR RETENTION_RATIO and ACTIVE_PRODUCERS are all replaced.

5.4.3 Feature Engineering

Three new features were obtained from existing data. We operationalised AGENCY_AGE as the difference between the year of record year and the year of agency appointment, which is the tenure of each agency. PREMIUM_GROWTH_YOY is the written premium percentage change year over year. It was capped between -100% and +1000% to avoid extreme outliers from skewing the model. LOG_WRTN_PREM_AMT was calculated on the basis of the log1p transformation of the raw premium target variable to normalise its distribution.

5.4.4 Categorical Encoding

Categorical variables PRODLINE and STATEABBR were transformed into numerical format. PRODLINE was labeled to become PRODLINEENC (0 = Commercial Lines, 1 = Personal Lines). STATEENC was derived from STATE_ABBR where 0 = Kentucky, 1 = Indiana, 2 = Michigan, 3 == Ohio, 4 == Pennsylvania, and lastly,, 5 = West Virginia.)

5.4.5 Train-Test Split

Cleaned data(147760 records) was split into 80% training and 20% test set with randomstate = 42 by using Scikit-learn traintest_split. There are 118208 records for Training set and 29552 Records for Test set.

5.5 Predictive Model Development

5.5.1 Model 1: Linear Regression

The baseline model implemented was a Linear Regression by using Scikit-learn LinearRegression class. The model only train a log-transformed target value(LOGWRTNPREM_AMT) with 13 input features. The snippet for the code is as below:

```
# Model 1: Linear Regression
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)

# Predicting
lr_pred = lr_model.predict(X_test)
```

Model 1: Linear Regression

Linear Regression was performed as expected for a baseline model with an R of 0.309. Relatively poor performance implies the correlation between the financial metrics and the volume of premium can not be linear, thus can't be modeled well by a linear regression model.

5.5.2 Model 2: Random Forest Regressor

The model that was implemented was Random Forest Regressor with 100 trees(`n_estimators = 100`) using parallel processing (`njobs=-1`) over all available cpu. Random Forest is an ensemble learning method which works by constructing multiple decision trees during training and outputs the mode of the classes (classification) or mean prediction (regression) of the individual trees. The snippet for the code is given below:

```
[32]: # Model 2: Random Forest Regressor
print("Training Random Forest... please wait")
rf_model = RandomForestRegressor(n_estimators=100, random_state=42, n_jobs=-1)
rf_model.fit(X_train, y_train)

# Predict
rf_pred = rf_model.predict(X_test)
```

Model 2: Random Forest Regressor




Figure 5.7: Rf Predicted vs Actual

5.5.3 Model 3: Facebook Prophet

Facebook Prophet was applied on yearly total written premium over time from 2005 to 2014. It was trained for the range of data from 2005 to 2012(8 years) and predict from 2013 to 2014(2 years). It required the column to be a date column(ds) and a numeric column(y). The changepointpriorscale parameter was used to make the trend somewhat flexible with a value of 0.05. The snippet for the code is given below:

```
# Model 3: Facebook Prophet
prophet_model = Prophet(yearly_seasonality=True,
                        changepoint_prior_scale=0.05)
prophet_model.fit(train_prophet)

# Forecast
future = prophet_model.make_future_dataframe(periods=2, freq='YS')
forecast = prophet_model.predict(future)
```

Model 3: Facebook Prophet

The negative R = -8.29 is due to a very small test data size of 2, hence it is impossible to get proper evaluation of R. As shown on the forecast visualization (figure 5.8), the Prophet model captured the trend with its plateau and a correctly-fitted confidence interval bound.




Figure 5.8: Prophet Forecasting




Figure 5.9: Model Comparison




Figure 5.10: Feature Importance

5.6 SQL Implementation

A SQLite database was created with name insurance.db with the processed data using sqlite3 package in Python and following 4 SQL queries were written:

Query 1: Annual Premium Summary

This query extracts the year, the written premium, average written premium per record and the new business premium from 2005 to 2014.

```

query1 = """
SELECT
    STAT_PROFILE_DATE_YEAR AS Year,
    COUNT(*) AS Total_Records,
    ROUND(SUM(WRTN_PREM_AMT), 2) AS Total_Written_Premium,
    ROUND(AVG(WRTN_PREM_AMT), 2) AS Avg_Premium_Per_Record,
    ROUND(SUM(NB_WRTN_PREM_AMT), 2) AS Total_New_Business_Premium
FROM insurance_sales
GROUP BY STAT_PROFILE_DATE_YEAR
ORDER BY STAT_PROFILE_DATE_YEAR
"""

result1 = pd.read_sql_query(query1, conn)
print("Query 1: Annual Premium Summary ===")
print(result1.to_string(index=False))
result1.to_csv('../data/processed/sql_q1_annual_summary.csv', index=False)

```

```

Query 1: Annual Premium Summary ===
Year  Total_Records  Total_Written_Premium  Avg_Premium_Per_Record  Total_New_Business_Premium
2005         12693         274847855.33             21653.50         33128543.64
2006         14059         413410568.17             29405.40         53804061.03
2007         14329         409016576.34             28544.67         44679899.65
2008         14908         415823965.06             27892.67         53784113.85
2009         15273         413354701.58             27064.41         50105748.60
2010         15092         410811002.20             27220.45         37370356.12
2011         14623         402377564.49             27516.76         27909844.77
2012         15639         405967783.45             25958.68         40334733.46
2013         15576         427281524.28             27432.04         58060760.09
2014         15568         437543915.53             28105.34         56464786.41

```

Query 2: Top 10 Agencies by Total Premium

This query was to list the best 10 agencies based on written premium in total and include average retention rate, and average loss ratio.

```

query2 = """
SELECT
    AGENCY_ID,
    ROUND(SUM(WRTN_PREM_AMT), 2) AS Total_Premium,
    ROUND(AVG(RETENTION_RATIO), 4) AS Avg_Retention,
    ROUND(AVG(LOSS_RATIO), 4) AS Avg_Loss_Ratio,
    COUNT(*) AS Years_Active
FROM insurance_sales
GROUP BY AGENCY_ID
ORDER BY Total_Premium DESC
LIMIT 10
"""

result2 = pd.read_sql_query(query2, conn)
print("Query 2: Top 10 Agencies by Total Premium")
print(result2.to_string(index=False))
result2.to_csv('../data/processed/sql_q2_top_agencies.csv', index=False)

```

Query 2: Top 10 Agencies by Total Premium					
AGENCY_ID	Total_Premium	Avg_Retention	Avg_Loss_Ratio	Years_Active	
5468	43800420.74	0.8477	0.8449	276	
9733	32454406.91	0.8609	0.3010	412	
1786	31314651.21	0.8776	0.4401	449	
886	29763095.55	0.8879	0.4722	338	
8365	27711251.47	0.8807	0.2287	355	
9362	27167460.59	0.8694	0.0019	494	
4351	24867838.69	0.8616	0.2347	443	
8652	24638042.63	0.8576	0.5744	361	
9761	24410979.89	0.8567	0.2429	273	
562	23899354.22	0.8862	0.2695	320	

Query 3: YoY Premium by Product Line

This query aims to investigate performance on product level by providing written premium, growth rate, customer retention rate of commercial and personal lines product in every year.

```

query3 = """
SELECT
    STAT_PROFILE_DATE_YEAR AS Year,
    PROD_LINE_ENC,
    ROUND(SUM(WRTN_PREM_AMT), 2) AS Total_Premium,
    ROUND(AVG(GROWTH_RATE_3YR), 4) AS Avg_3Yr_Growth_Rate,
    ROUND(AVG(RETENTION_RATIO), 4) AS Avg_Retention
FROM insurance_sales
GROUP BY STAT_PROFILE_DATE_YEAR, PROD_LINE_ENC
ORDER BY PROD_LINE_ENC, STAT_PROFILE_DATE_YEAR
"""

result3 = pd.read_sql_query(query3, conn)
print(" Query 3: YoY Premium by Product Line")
print(result3.to_string(index=False))
result3.to_csv('../data/processed/sql_q3_product_line_yoy.csv', index=False)

```

Query 3: YoY Premium by Product Line

Year	PROD_LINE_ENC	Total_Premium	Avg_3Yr_Growth_Rate	Avg_Retention
2005	0	95495332.21	-0.0156	0.8875
2006	0	151997178.86	-0.0156	0.8875
2007	0	152532681.59	-0.0156	0.8875
2008	0	165522754.49	0.0432	0.8875
2009	0	167900101.65	0.0413	0.8875
2010	0	166666352.74	0.0290	0.8875
2011	0	166544290.38	0.0291	0.8875
2012	0	174505752.49	0.0236	0.8875
2013	0	190573499.00	0.0342	0.8875
2014	0	204441129.84	0.0545	0.8875
2005	1	179352523.12	-0.0156	0.8325
2006	1	261413389.31	-0.0156	0.8288
2007	1	256483894.75	-0.0156	0.8425
2008	1	250301210.57	-0.0008	0.8519
2009	1	245454599.93	-0.0015	0.8609
2010	1	244144649.46	0.0102	0.8561
2011	1	235833274.11	0.0159	0.8472
2012	1	231462030.96	-0.0009	0.8491
2013	1	236708025.28	0.0222	0.8551
2014	1	233102785.69	0.0192	0.8434

Query 4 State-wise Financial Health

This query analyze the financial health on each state by showing the total written premium, loss ratio, customer retention rate, and new business premium by state.

```

query4 = """
SELECT
    STATE_ENC,
    COUNT(DISTINCT AGENCY_ID) AS Unique_Agencies,
    ROUND(SUM(WRTN_PREM_AMT), 2) AS Total_Premium,
    ROUND(AVG(LOSS_RATIO), 4) AS Avg_Loss_Ratio,
    ROUND(AVG(RETENTION_RATIO), 4) AS Avg_Retention_Ratio,
    ROUND(SUM(NB_WRTN_PREM_AMT), 2) AS Total_New_Business
FROM insurance_sales
GROUP BY STATE_ENC
ORDER BY Total_Premium DESC
"""

result4 = pd.read_sql_query(query4, conn)
print("Query 4: State-wise Financial Health")
print(result4.to_string(index=False))
result4.to_csv('../data/processed/sql_q4_state_health.csv', index=False)

```

STATE_ENC	Unique_Agencies	Total_Premium	Avg_Loss_Ratio	Avg_Retention_Ratio	Total_New_Business
3	832	2.351810e+09	0.5282	0.8682	219158272.40
4	390	5.280010e+08	0.7641	0.8802	70329301.52
1	363	4.690770e+08	0.6287	0.8680	74695836.54
0	444	4.388524e+08	0.7583	0.8594	56222247.30
5	185	1.824008e+08	0.8823	0.8761	21054516.55
2	108	4.029395e+07	0.8632	0.8875	14182673.31

5.7 Power BI Dashboard

An interactive dashboard was created in Power BI Desktop with 3 pages using the processed data and the SQL queries. The dashboards cover descriptive, financial and prediction aspects:

Page 1: Sales Overview Dashboard

This page presents an executive overview to help identify the sales trend with 4 KPI indicators (Total Written Premium: \$4.01bn, Total Policy Records: 147.76k, Total New Business Premium: \$455.64m, Average Loss Ratio), one line chart showing annual premium trend, one bar chart of top 10 agencies and one cluster column chart comparing new business premium and total premium. Slicers by Year range and product line were provided for the user to dynamically analyze data.

Page 2: Financial Health Dashboard

This page shows the financial indicators with one line chart showing the difference between Commercial and Personal lines performance, one gauge to show the average customer

retention rate(86.9%), one treemap representing written premium by state and one line chart showing average loss ratio. These visuals answer the need of monitoring the financial health of the insurance portfolio.

Page 3: Predictive Analytics Dashboard

This page presents the model results with a Scatter chart comparing predicted value of the Random Forest model with actual value and a Sales Forecast chart showing the forecast of annual total written premium from the Prophet model with the prediction interval range, as well as the comparison between R scores and MAE scores of the two models using Bar chart.

CHAPTER 6: TESTING

6.1 Testing Strategy

This project's testing is structured in three levels. Each level is based on a different type of software testing that aims at fulfilling specific requirements. Tests were done manually:

- Unit Testing
- Integration Testing
- System Testing

6.2 Unit Testing

TC-ID	Component	Input	Expected Output	Status
UT-01	Data Loading	finalapi.csv path	DataFrame shape (213328, 49)	PASS
UT-02	Null Check	Loaded DataFrame	0 null values in all columns	PASS
UT-03	Duplicate Check	Loaded DataFrame	0 duplicate rows	PASS
UT-04	Year Filter	Remove year 2015	Shape reduced to 192,347	PASS
UT-05	Negative Filter	Remove WRTN_PREM_AMT ≤ 0	Final shape 147,760	PASS
UT-06	Sentinel Replace	99999 in RETENTION_RATIO	Replaced with median 0.8875	PASS
UT-07	Log Transform	WRTN_PREM_AMT values	LOG_WRTN_PREM_ AMT column created	PASS
UT-08	Train-Test Split	147,760 records, 80/20	118,208 train, 29,552 test	PASS

UT-09	LR Model Train	X_train, y_train	Model fitted, no errors	PASS
UT-10	RF Model Train	X_train, y_train, n=100	Model fitted, no errors	PASS
UT-11	Prophet Format	Annual aggregated data	ds and y columns created	PASS
UT-12	MAE Calculation	y_test, lr_pred	MAE = 1.470	PASS

6.3 Integration Testing

TC-ID	Integration Point	Expected Behavior	Actual Behavior	Status
IT-01	Data Load → Preprocess	Cleaned data saved to processed/	insurance_cleaned.csv created	PASS
IT-02	Preprocess → Model Train	X_train/y_train loaded correctly	Models trained successfully	PASS
IT-03	Model → Evaluation	Metrics computed for all models	MAE, RMSE, R ² all returned	PASS
IT-04	Data → SQLite	CSV loaded into database	147,760 rows in table	PASS
IT-05	SQL → CSV Export	Query results saved as CSV	4 CSV files created	PASS
IT-06	CSV → Power BI	Files imported into dashboard	3 tables loaded correctly	PASS
IT-07	Model → Power BI	Predictions exported to CSV	rf_predictions.csv, prophet_forecast.csv created	PASS

6.4 System Testing

TC-ID	System Test	Expected Behavior	Actual Behavior	Status
ST-01	Full Pipeline Run	All 3 notebooks run without errors	Completed successfully	PASS
ST-02	Model Comparison	RF outperforms LR on R^2	RF: 0.848 > LR: 0.309	PASS
ST-03	Dashboard Page 1	KPI cards display correct totals	\$4.01bn total premium shown	PASS
ST-04	Dashboard Slicer	Year slicer filters all visuals	All charts update on filter	PASS
ST-05	Dashboard Page 3	Scatter plot shows positive correlation	Clear diagonal pattern observed	PASS
ST-06	GitHub Push	All files pushed to repository	Repository updated successfully	PASS

CHAPTER 7: RESULTS AND DISCUSSION

7.1 EDA Results

The exploratory data analysis revealed that the total written premium over 10 years (2005-2014) was \$4.01B for the insurance agency portfolio and there were 147,760 records from 1,623 agencies. The portfolio was found to have had a rapid growth period (50.5% increase in written premium) during 2005-2006 and since then has plateaued at around \$400-440M in annual written premium, and in 2014 this has reached \$437.5M.

Ohio was clearly the predominant state in the portfolio contributing \$2,450.3M (58%) of the total written premium, and is significantly ahead of its nearest rivals Pennsylvania (\$554.9M), Kentucky (\$491.8M), and Indiana (\$454.7M). Customer retention ratio throughout the portfolio was 88.75% and this suggests loyalty to the insurance agencies in the portfolio. The loss ratio was volatile throughout the portfolio with a spike in 2012 possibly suggesting a high claims period.

Commercial Lines showed an increasing trend throughout the period (+114% 2005-2014) and whereas the Written Premium values from Personal Lines was high, this had decreased by the end of 2014 after peaking in 2006.

7.2 Model Performance Results

Model	MAE	RMSE	R ² Score
Linear Regression	1.470	1.830	0.309
Random Forest	0.505	0.858	0.848
Facebook Prophet	\$14,416,640	\$15,642,871	-8.294 (2 test points)




Figure 7.1: Model Comparison

The Random Forest Regressor performs well, yielding an R score of 0.848. This suggests the model explains 84.8% of the variation of the log written premium values. This is an improvement of 174% compared to the Linear Regression model, which was measured with an R score of 0.309. MAE on the log scale is 0.505, suggesting that on average, our prediction is less than 10% from the actual values.

Our analysis into the importance of features for prediction (see Figure 5.10) indicates that PRDINCRDLOSSESAMT (incurred losses) is our most important feature, importance of 0.35, followed by NBWRTNPREMAMT (new business premium, importance 0.26) and PREMIUMGROWTHYOY (year-over-year premium growth, importance 0.15). The practical implication of this is that for a firm attempting to forecast their sales most accurately they will need to focus on providing correct incurred losses figures and have the right management system to facilitate strong growth of new business.

7.3 Dashboard Results




Figure 7.2: Sales Overview Dashboard

The Sales Overview Dashboard (page 1) of the Power BI report can be seen in Figure 7.2. The key performance indicator cards give an overall picture of the insurance agency portfolio by stating a total written premium figure of \$4.01 Billion, 147,760 policy records, and new business premium of \$455.64 Million. The bar chart clearly shows the 2005/2006 peak and plateau, and the top 10 agency bar chart clearly illustrates the dominant position of Agency 5468 with \$44 Million of premium followed by Agency 9733 (\$32 Million) and Agency 1786 (\$31 Million).




Figure 7.3: Financial Health Dashboard

Figure 7.3: Financial Health DashboardThe Financial Health Dashboard (page 2) is shown in Figure 7.3. The retention rate gauge illustrates a strong customer loyalty across the portfolio with a reading of 86.9%. The treemap showing premium broken down by state highlights Ohio at \$2.35 Billion as clearly the most dominant state in the portfolio, and is easily visible. The loss ratio trend chart for the portfolio clearly illustrates the 2012 spike to 17.3K (this was due to the sentinel value residual) and recovery by 2014 to 6.5K.




Figure 7.4: Predictive Analytics Dashboard

Figure 7.4: Predictive Analytics DashboardFigure 7.4 displays the Predictive Analytics Dashboard (page 3). The Random Forest scatter plot shows a good fit between the predicted and actual values with the points forming a tight cluster around the perfect prediction line, and this is true across the whole range of premium figures. The Prophet forecasting chart show how the trend is correctly estimated with confidence intervals in place. The comparative charts again highlight how the Random Forest prediction model is superior for this forecasting problem to a standard linear regression.

7.4 Comparison with Existing Systems

Existing manual forecasting methods in the regional insurance market and simple statistical methods are significantly inferior to the system developed here. Random Forest regressor yields an R score of 0.848, while manual regression models only tend to achieve a result between 0.3-0.5 according to a study into sales forecasting. Furthermore, compared to traditional methods, feature importance rankings are obtained with our model, providing insight into key contributors to premium.

CHAPTER 8: CONCLUSION AND FUTURE ENHANCEMENTS

8.1 Conclusion

In this project we have shown the capability of using predictive analytics for forecasting insurance sales in finance industry and we have fulfilled all 4 objectives of the project. According to the results obtained from evaluation metrics, Random Forest Regressor is the best model with $R=0.848$ and $MAE=0.505$. It proves that ensemble machine learning methods significantly out-perform the linear model when it is used to predict insurance premiums. We studied the 147,760 clean records of 1,623 insurance agencies for a 10-year period (2005-2014) and developed and compared 3 different predictive models to provide insights through a three-page interactive Power BI dashboard to insurance agencies.

We fulfilled our first objective of studying the importance of predictive analytics in finance by conducting a literature review combining data from 5 peer-reviewed papers and identifying a research gap in agency-level insurance premium forecasting. We satisfied our second objective of developing predictive models through Linear Regression, Random Forest Regressor and Prophet; the models were fed by a pipeline designed to preprocess data to remove sentinel values, engineer new features and normalize the target variable.

To fulfill our third objective, which was to compare the accuracy of the developed predictive models, we utilized three performance metrics: MAE, RMSE and R. The Random Forest Regressor was found to be the model that performed best with $R = 0.848$, $MAE = 0.505$, proving that ensemble machine learning approaches perform significantly better than linear models when forecasting insurance premiums. We fulfilled our fourth and final objective of providing business recommendations by examining the feature importance ranks, the spatial concentrations of agencies, and the trends of various products.

Among the technical achievements of the project, there is: the achievement of a clean dataset with zero-null and zero-duplicates; the 174% improvement from baseline model to best model regarding R; the detection that incurred losses and new business premium are the two main drivers of insurance sales; and a fully interactive three-page Power BI dashboard with both descriptive and predictive data.

8.2 Future Scope

Further improvements of the developed predictive analytics system could be achieved through:

- **Hyperparameter Tuning:** Implementing a search strategy using Grid Search or Random Search cross-validation would allow fine-tuning of Random Forest parameters such as `n_estimators`, `max_depth`, and `min_samples_split` to improve the model beyond $R = 0.848$.
- **Advanced Models:** implementing other ensemble methods like XGBoost or LightGBM could achieve higher accuracy due to their use of gradient boosting which is very suitable for tabular financial data.
- **Real-Time Data Integration:** linking the forecasting pipeline to a live agency management system through an API could provide instantaneous prediction updates instead of the current batch processing.
- **Expanded Geographic Coverage:** analyzing additional states and regions instead of limiting to the studied states would reduce spatial concentration risk and make the model generalizable.
- **Deep Learning Integration:** employing LSTM neural networks for time series prediction could better capture longer term temporal patterns in premium data which are undetectable by Prophet.
- **Cloud Deployment:** deploying the pipeline in an enterprise platform like AWS SageMaker or Azure Machine Learning would provide scalability and a production-ready system to serve predictions to enterprise insurance clients.
- **Mobile Dashboard:** creating a Power BI dashboard optimized for mobile devices would allow field agents and regional managers to view predictions while in transit.

CHAPTER 9: REFERENCES

Balasubramanian, Ramnath, Libarikian, Anand, & McElhaney, Doug. (2021). Insurance 2030: How AI will reshape the insurance industry.

McKinsey and Company.

<https://www.mckinsey.com/industries/financial-services/our-insights/insurance-2030>

Boustani, Nadia, Emrouznejad, Ali, Gholami, Roya, Despic, Ozren, & Ioannou, Alexis. (2023). Accurate consumer loan cross-selling prediction using deep learning networks. *Annals of Operations Research*, 339(1-2), 613-630.

<https://doi.org/10.1007/s10479-023-05209-5>

Breiman, Leo. (2001). Random forests. *Machine Learning*, 45(1), 5-32.

<https://doi.org/10.1023/A:1010933404324>

Tianqi Chen & Carlos Guestrin (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785-794.

[XGBoost | Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining](#)

Davenport, Thomas H., & Harris, Jeanne G. (2007). *Competing on analytics: The new science of winning*. Harvard Business School Press. [Competing on Analytics: The New Science of Winning | Guide books | ACM Digital Library](#)

Du, Shengdong, Hao, Dong, & Li, Xin. (2022). Stock price forecasting research using the random forest algorithm. *IEEE International Conference on Data Science and Computer Application (ICDSCA)*. <https://doi.org/10.1109/ICDSCA56264.2022.9987903>

Eckert, Christian, Neunsinger, Christina, & Osterrieder, Kristina. (2022). Customer satisfaction management: Digital tools for insurance companies. *The Geneva Papers on Risk and Insurance - Issues and Practice*, 47(3), 569-602. [Forecasting, Structural Time Series Models and the Kalman Filter](#)

Elalem, Yacine Kouame, Maier, Stefan, & Seifert, Ralf W. (2023). Machine learning-based framework for new product sales forecasting with short life cycles using deep neural

networks. International Journal of Forecasting, 39(4), 1874-1894.
<https://doi.org/10.1016/j.ijforecast.2022.07.003>

Harvey, Andrew C. (1990). Forecasting, structural time series models, and the Kalman filter. Cambridge University Press.

[Forecasting, Structural Time Series Models and the Kalman Filter](#)

Jiang, Ming, & Wang, Xiaolong. (2021). Intelligent prediction method for listed enterprises' financial crisis using the random forest algorithm. Mathematical Problems in Engineering. <https://doi.org/10.1155/2021/3807480>

Kaggle. (2023). Insurance agency performance dataset - moneystore/agencyperformance. Kaggle Platform. <https://www.kaggle.com/datasets/moneystore/agencyperformance>

Kaushik, Keshav, Bhardwaj, Abhishek, Dwivedi, Ashutosh Dhar, & Singh, Rajani. (2022). Machine learning-based regression framework to predict health insurance premiums. International Journal of Environmental Research and Public Health, 19(13), 7898.
<https://doi.org/10.3390/ijerph19137898>

Kumar, Nikhil. (2023). A study of the impact and applications of predictive analytics in sales forecasting. International Journal for Research in Applied Science and Engineering Technology, 11(12). <https://doi.org/10.22214/ijraset.2023.57535>

Nwaimo, Chukwuemeka Sunday, Adegbola, Adedayo Emmanuel, & Adegbola, Modupe Dorcas. (2024). Predictive analytics for financial decision-making: Trends, challenges and opportunities. Global Journal of Research in Multidisciplinary Studies, 2, 16-26.
<https://gjrms.com/index.php/home/article/view/52>

Olagoke, Michael Femi. (2025). The role of predictive analytics in enhancing financial performance. Journal of Financial Risk Management. [The Role of Predictive Analytics in Enhancing Financial Decision-Making and Risk Management](#)

Pedregosa, Fabian, Varoquaux, Gael, Gramfort, Alexandre, Michel, Vincent, Thirion, Bertrand, Grisel, Olivier, Blondel, Mathieu, Prettenhofer, Peter, Weiss, Ron, Dubourg, Vincent, Vanderplas, Jake, Passos, Alexandre, Cournapeau, David, Brucher, Matthieu, Perrot, Matthieu, & Duchesnay, Edouard. (2011). Scikit-learn: Machine learning in

Python. Journal of Machine Learning Research, 12, 2825-2830. [Scikit-learn: Machine Learning in Python](#)

Prakash, Vijaya Sundar, Bushra, Syed Nawaz, Subramanian, Nithyanandam, Indumathy, D., Mary, Selva Arul Latha, & Thiagarajan, Rajendran. (2022). Random forest regression with hyper parameter tuning for medical insurance premium prediction. International Journal of Health Sciences, 6(S6), 7093-7101. . <https://doi.org/10.53730/ijhs.v6nS6.10993>

Taylor, Sean J., & Letham, Benjamin. (2018). Forecasting at scale. The American Statistician, 72(1), 37-45. <https://doi.org/10.1080/00031305.2017.1380080>

Vellela, Siva Sankara, Pushpalatha, D., Sarathkumar, G., Kavitha, C. H., & Harshithkumar, D. (2023). Advanced intelligence health insurance cost prediction using random forest. SSRN Working Paper. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4473700

World Economic Forum. (2023). The future of jobs report 2023. World Economic Forum. [The Future of Jobs Report 2023 | World Economic Forum](#)

CHAPTER 10: APPENDICES

Appendix A: Project Repository

The complete project source code, Jupyter notebooks, SQL queries, and dashboard files are available on GitHub :

[Insurance Sales Forecasting](#)

Appendix B: Dataset Information

Dataset: Insurance Agency Performance Dataset

Source: [Insurance Data](#) - Kaggle

License: Open Data / Public Domain

Original Records: 213,328 | Cleaned Records: 147,760 | Features: 49

Appendix C: Installation Guide

Follow these steps to create a local copy of an existing remote repository. :

1. Clone the repository: git clone [Insurance Sales Forecasting](#)
2. Install dependencies: pip install pandas numpy matplotlib seaborn scikit-learn prophet
3. Download the dataset from [Kaggle](#) and place in data/raw/ directory
4. Run notebooks in order: 01_data_exploration.ipynb, 02_preprocessing_modeling.ipynb, 03_sql_queries.ipynb
5. Open insurance_dashboard.pbix in Power BI Desktop and refresh data sources

Appendix D: Model Performance Summary

Model	MAE	RMSE	R ²	Scale
Linear Regression	1.470	1.830	0.309	Log
Random Forest	0.505	0.858	0.848	Log
Facebook Prophet	14,416,640	15,642,871	-8.294	USD